


Operating Systems

04. Threads

Prof. Dr. Frank Bellosa | WT 2020/2021

KARLSRUHE INSTITUTE OF TECHNOLOGY (KIT) – ITEC – OPERATING SYSTEMS

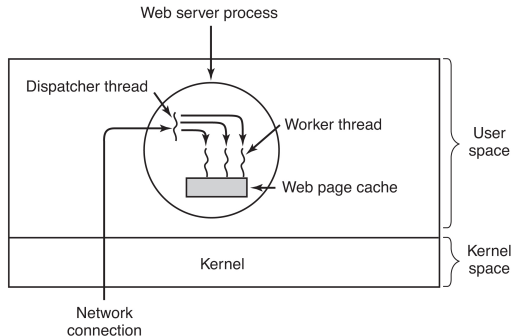


Processes vs. Threads

- In traditional OSeS, each process has
 - it's own address space (AS)
 - it's own set of allocated resources
 - one thread of execution (one execution state)
- Modern OSeS handle processes and threads of execution more flexibly
 - Processes provide the abstraction of an address space and resources
 - Threads provide the abstraction for execution states of that AS
- Research OSeS already offer multiple address spaces to different threads of the same process
 - e.g. for live code patching of multi-threaded processes
see related work [RDF⁺20]


Why have multiple threads at all?

- Many programs do multiple “things” at once
- Multi-threaded web server [TB15]
 - Accept new connections
 - Read request from client
 - Fetch required data
 - Process and deliver data
- Multiple threads can potentially run in parallel on a multiprocessor
- Some of these activities may **block**
- Writing a program made of many sequential threads may be easier than avoiding blocking operations (see related work “Threads vs. Events”)





Why have multiple threads per process?

- Example: Lots of shared state and blocking operations [TB15]
 - Word processor: Read input, format output, write backup file



Multiples Processes vs. Multiple Threads

- Whether to use multiple processes or threads depends on the activity [TB15]
- Processes rarely share data, and if they do, they do it explicitly.
- Closely related activities share data which favors threads.



Thread Libraries

- Thread libraries provide an API for creating and managing threads
- Pthreads
 - POSIX API for thread management and synchronization (IEEE 1003.1c)
 - API specifies behavior of the thread library (> 60 API calls)
 - Internal details are up to the specific implementation of the library
 - Common in UNIX operating systems (Linux, Mac OS X, Solaris, AIX)

Basic POSIX Thread API

- Each `pthread` is associated with
 - an identifier (Thread ID, TID)
 - a set of registers (including IP and SP)
 - a stack area to hold the execution state (functions/local vars) of that thread
- `Pthread_create` Create a new thread
 - Pass: pointer to `pthread_t` (will hold TID after successful call)
 - Pass: attributes, start function, and arguments
 - Returns: 0 on success or error value
- `Pthread_exit` Terminate the calling thread
 - Pass: exit code (casted to a void pointer)
 - Free's resources (e.g., stack)
- `Pthread_join` Wait for a specific thread to exit
 - Pass: `pthread_t` to wait for (or -1 for any thread)
 - Pass: Pointer to pointer for exit code
 - Returns: 0 on success, otherwise error value
- `Pthread_yield` Release the CPU to let another thread run

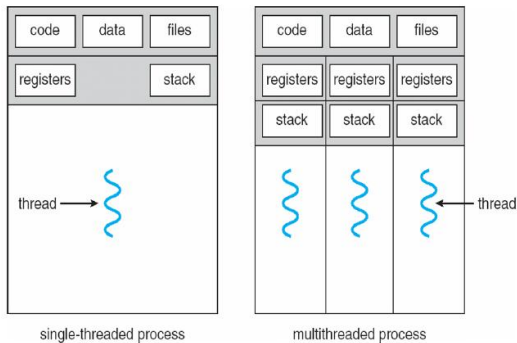
Pthread Example

```
void* greet( void *id )
{
    printf( "Hello, I am %ld\n", (intptr_t) id );
    pthread_exit( (void*) 0 );
}

int main()
{
    pthread_t threads[NUM];
    for( int i = 0; i < NUM; ++i )
    {
        int status = pthread_create( threads + i, NULL,
                                    greet, (void *) (intptr_t) i );
        if( status != 0 )
            die( "Error creating thread" );
    }
    for( int i = 0; i < NUM; ++i )
        pthread_join( threads[i], NULL );
    return 0;
}
```


Data Structures for Processes and Threads

- Processes group resources
 - Threads encapsulate execution
- Each of those abstractions requires different data



[SGG12]

PCB vs. TCB

- We differentiate between
 - **Process Control Block (PCB)**: Information needed to implement processes
 - **Thread Control Block (TCB)**: Per thread data

- Typical items in each category are:

PCB	TCB
Address space	Instruction pointer
Open files	Registers
Child processes	Stack
Pending alarms	State


- The PCB is always known to the OS
- Whether or not the OS kernel knows about threads or not depends on the **thread model**

Thread Model Overview

- The OS kernel always knows of at least one thread per process
 - Threads that are known to the OS kernel are called **kernel threads**
 - Threads that are known to the process are called **user threads**
- Threads can be fully implemented in user-space
 - **Many-to-One Model**: The kernel only knows one of possibly multiple threads
 - User threads in this model are called **User Level Threads (ULT)**
- The kernel can be fully aware of and responsible for managing threads
 - **One-to-One Model**: Each user thread maps to a kernel thread
 - User threads in this model are called **Kernel Level Threads (KLT)**
- The kernel can know of multiple threads per process, yet there are even more threads known to the process
 - **M-to-N Model**: Flexible mapping of user threads to less kernel threads
 - Also known as **hybrid thread model**

Many-to-One Model: User Level Threads (ULT)

- Kernel only manages process → multiple threads unknown to kernel
- Threads managed in user-space library (e.g., GNU Portable Threads)
- + Faster thread management operations (up to 100 times)
- + Flexible scheduling policy
- + Few system resources
- + Can be used even if the OS does not support threads
- No parallel execution
- Whole process blocks if only one user thread blocks
- Need to re-implement parts of the OS (e.g., scheduler)[TB15]



ULT Implementation

- Unix like systems (e.g., Linux) define
 - the types `mcontext_t` and `ucontext_t` to keep thread state
 - `makecontext` Initialize a new context
 - `getcontext` Store currently active context
 - `setcontext` Replace current context with different one
 - `swapcontext` User-level context switching between threads
- Using those functions, calls for creating threads, yielding, wait, can easily be implemented fully in user-space
 - e.g., `yield` saves own context and replaces itself with a different context
- Periodic thread switching can be implemented using a `SIGALRM` signal handler
- We will distribute an example how to use these function with an assignment in the tutorials (`ult.h/ult.c`)
- An alternative interface to the context would be `setjmp` and `longjmp`

Address-Space Layout with Two User Level Threads

■ Stack


- “Main stack” known to OS is used by thread library (e.g., called via **SIGALRM** **upcalls**)
- Own execution state (\equiv stack) for every thread is allocated dynamically by user thread library on the heap using **malloc**

■ Heap

- Concurrent heap use possible
- Attention: not all heaps are reentrant!

■ Data


- Is divided into BSS, data and read-only data



One-to-One Model: Kernel Level Threads (KLT)

- Kernel knows and manages every thread
 - Every thread known by kernel maps to one thread known by user
 - Windows, Linux, Open Solaris, Mac OS X all support this

- + Real parallelism possible
 - + Threads block individually
 - OS manages every thread in the system (TCB, stacks, ...)
 - Syscalls needed for thread management
 - Scheduling fixed in OS
- [TB15]



Address-Space Layout with Two Kernel Level Threads

■ Stack


- Own execution state (\equiv stack) for every thread

■ Heap

- Parallel heap use possible
- Attention: not all heaps are thread-safe!
- Even if the heap is thread-safe:
Not all heap implementations perform well with many threads

■ Data

- Is divided into BSS, data and read-only data



KLT Implementation and Issues

- All thread management data is stored in the kernel
- Thread management functions are provided as syscalls
- What happens when a process with multiple KLTs calls `fork`?
- **Signals** are used in UNIX systems to notify a process that a particular event has occurred
 - e.g., process can ask OS to send **SIGALRM** after a specific time
 - **signal handler** can run
 - on the process stack
 - on a stack, dedicated to the specific signal handler
 - on a stack, dedicated to all signal handlers
 - Whom is the signal delivered to?
 - All threads in process?
 - One thread that receives all signals?
 - The thread that set up the handler? What if multiple threads subscribe this signal?

M-to-N Model: Hybrid Threads


- **M** ULTs are mapped to (at most) **N** KLTs [TB15]

- Goal: pros of ULT and KLT – non-blocking with quick management
- Create a “sufficient” number of KLTs (kernel is only aware of KLTs)
- Flexibly allocate ULTs on those KLTs

- + Flexible scheduling policy
- + Efficient execution

- Hard to debug
- Hard to implement (blocking, variable number of KLTs)

e.g., Solaris 9 and earlier, Windows NT/2000 with ThreadFiber, Linux 2.4 [MSLM91][KAR⁺06]



Hybrid Thread Implementation: Scheduler Activations

- Goal: Don't involve kernel on thread activities such as `create` and `join`
- Idea: Map multiple ULTs on each KLT
 - When a ULT blocks (e.g., page fault, syscall) the user-space run-time system can dispatch a different ULT after being notified by the kernel
- Approach: [Upcalls](#)
 - The kernel notices that a thread will block and sends an upcall to the process
 - The upcall notifies the process of the thread id and event that happened
 - The upcall handler of the process can then schedule a different thread in that process
 - The kernel later informs the process that the blocking event has finished via another upcall

[ABLL91]

Making Single-Threaded Code Multithreaded

- It is hard to make single-threaded code multithreaded
- Not all state should be shared between threads
 - `errno` contains the error number of the last syscall (0 on no error)
 - `errno` is overwritten on subsequent system calls
 - Which thread does the current value belong to?
- Much existing code, including many libraries, are not re-entrant
 - `malloc` is not always thread-safe
 - `strtok` is not thread-safe (use `strtok_r`)
 - Generally: use `_r` variants of functions (`rand_r` instead of `rand`)
- How should stack growth be managed?
 - Normally the kernel grows the (single) stack automatically when needed
 - What if there are multiple stacks?

References I

- [ABLL91] Thomas E. Anderson, Brian N. Bershad, Edward D. Lazowska, and Henry M. Levy. Scheduler activations: Effective kernel support for the user-level management of parallelism. *SIGOPS Oper. Syst. Rev.*, 25(5):95–109, September 1991.
- [KAR⁺06] Orran Krieger, Marc Auslander, Bryan Rosenburg, Robert W. Wisniewski, Jimi Xenidis, Dilma Da Silva, Michal Ostrowski, Jonathan Appavoo, Maria Butrico, Mark Mergen, Amos Waterland, and Volkmar Uhlig. K42: Building a complete operating system. In *Proceedings of the 1st ACM SIGOPS/EuroSys European Conference on Computer Systems 2006*, EuroSys '06, page 133–145, New York, NY, USA, 2006. Association for Computing Machinery.
- [MSLM91] Brian D. Marsh, Michael L. Scott, Thomas J. LeBlanc, and Evangelos P. Markatos. First-class user-level threads. *SIGOPS Oper. Syst. Rev.*, 25(5):110–121, September 1991.

References II

- [RDF⁺20] Florian Rommel, Christian Dietrich, Daniel Friesel, Marcel Köppen, Christoph Borchert, Michael Müller, Olaf Spinczyk, and Daniel Lohmann. From global to local quiescence: Wait-free code patching of multi-threaded processes. In *14th Symposium on Operating System Design and Implementation (OSDI '20)*, pages 651–666, November 2020.
- [SGG12] Abraham Silberschatz, Peter B. Galvin, and Greg Gagne. *Operating System Concepts*. Wiley Publishing, 9th edition, 2012.
- [TB15] Andrew S Tanenbaum and Herbert Bos. *Modern operating systems*. Pearson, 4th edition, 2015.