


# Operating Systems

## 08. Page Fault Handling

Prof. Dr. Frank Bellosa | WT 2020/2021

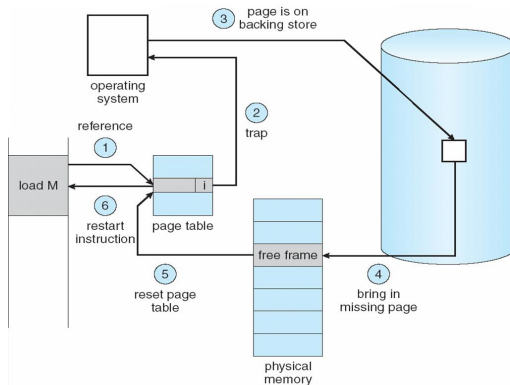
KARLSRUHE INSTITUTE OF TECHNOLOGY (KIT) – ITEC – OPERATING SYSTEMS



# Page Faults

# Page Fault Handling [KELS62]

- Access to page that is currently not present in main memory causes page fault (exception that invokes OS)
- OS checks validity of access (requires additional info)
- Get empty frame (free or victim)
- Prepare requested page (e.g., clean or load content from disk into frame)
- Adapt page table
- Set present bit of respective entry (a.k.a. as setting **valid-invalid bit** to **v**)
- Restart instruction that caused the page fault [SGG12]



# Page Replacement [SGG12]

- Save/clear victim page (see chapter on page replacement policies)
  - Drop page if fetched from disk (e.g., code) and clean (PTE dirty bit)
  - Write back modifications if from disk and dirty
  - Write pagefile/swap partition otherwise (e.g., stack, heap memory)

- Unmap page from old AS


- unset valid bit in PTE
- flush TLB

- Prepare the new page

- e.g., NULL page
- e.g., load new contents

- Map the page frame into the new address space(s)

- set valid bit in PTE
- flush TLB



# Page Replacement [SGG12]

- Save/clear victim page (see chapter on page replacement policies)
  - Drop page if fetched from disk (e.g., code) and clean (PTE dirty bit)
  - Write back modifications if from disk and dirty
  - Write pagefile/swap partition otherwise (e.g., stack, heap memory)

- Unmap page from old AS


- unset valid bit in PTE
  - flush TLB

- Prepare the new page

- e.g., NULL page
  - e.g., load new contents

- Map the page frame into the new address space(s)

- set valid bit in PTE
  - flush TLB



# Page Replacement [SGG12]

- Save/clear victim page (see chapter on page replacement policies)
  - Drop page if fetched from disk (e.g., code) and clean (PTE dirty bit)
  - Write back modifications if from disk and dirty
  - Write pagefile/swap partition otherwise (e.g., stack, heap memory)

- Unmap page from old AS



- unset valid bit in PTE
- flush TLB

- Prepare the new page

- e.g., NULL page
- e.g., load new contents


- Map the page frame into the new address space(s)

- set valid bit in PTE
- flush TLB



# Page Replacement [SGG12]

- Save/clear victim page (see chapter on page replacement policies)
  - Drop page if fetched from disk (e.g., code) and clean (PTE dirty bit)
  - Write back modifications if from disk and dirty
  - Write pagefile/swap partition otherwise (e.g., stack, heap memory)
- Unmap page from old AS
  - unset valid bit in PTE
  - flush TLB
- Prepare the new page
  - e.g., NULL page
  - e.g., load new contents
- Map the page frame into the new address space(s)
  - set valid bit in PTE
  - flush TLB



# Page Fault Latency

- Page Fault Rate  $0 \leq p \leq 1.0$ 
  - $p = 0$ : No page faults
  - $p = 1$ : Every reference is a page fault
- Effective Access Time (EAT)

$$\text{EAT} = (1 - p) \times \text{memory access} + p \times \left( \begin{array}{l} \text{page fault overhead} \\ + \text{page fault service time} \\ + \text{restart overhead} \end{array} \right)$$



# Performance Impact of Page Faults

- Memory access time = 200 nanoseconds
- Average page fault service time = 8 milliseconds

■

$$\begin{aligned} \text{EAT} &= (1 - p) \times 200 + p(8\text{ms}) \\ &= (1 - p) \times 200 + p \times 8,000,000 \\ &= 200 + p \times 7,999,800 \end{aligned}$$

- If one access out of 1,000 causes a page fault, then EAT = 8.2 microseconds.  $\Rightarrow$   
Slowdown by a factor of 40!

# What to fetch?

- Bring in page that caused page fault
- Pre-fetch surrounding pages?
  - Reading two disk blocks is approximately as fast as reading one
  - As long as no track/head switch, seek time dominates (disk)
  - If application exhibits spatial locality → big win
- Pre-zero pages?
  - Don't want to leak information between processes
  - Need 0-filled pages (**0-pages**) for stack, heap, .bss, ...
  - Zero on demand?
  - Keep a pool of 0-pages that is filled in the background when the CPU is idle?

# How to resume a process after a fault?


- Hardware provides info about page fault
  - Faulting virtual address: `%cr2` on intel
- OS needs to figure out context of fault. Was the instruction a
  - Read or write?
  - Instruction fetch?
  - User access to kernel memory?
- Idempotent instructions are easy
  - Just re-do load/store instructions
  - Just re-execute instructions that only access one address

# Complex instructions must be re-started, too


- Some CISC instructions are difficult to restart such as
  - Block move of overlapping areas, string move instructions
  - Auto-increments/decrements of multiple locations
  - Instructions that keep and update state in source `%esi`, destination `%edi`, and counter `%ecx` registers
- Possible Solutions
  - Touch all relevant pages before operation starts
  - Keep modified data in registers so that page faults can't take place
  - Design ISA such that complex operations can execute partially and leave consistent state on a page fault (easy job for the OS)

# Shared Library/Code Using Virtual Memory

## [SGG12]



# Memory-Mapped Files [SGG12]




## Other Issues – Memory-Mapped Files

- Memory-mapped file I/O allows file I/O to be treated as routine memory access by **mapping** a disk block to a page in memory
- A file is initially read using demand paging. A page-sized portion of the file is read from the file system into a physical page. Subsequent reads/writes to/from the file are treated as ordinary memory accesses
- Simplifies file access by treating file I/O through memory rather than `read()` `write()` system calls
- Also allows several processes to map the same file allowing the pages in memory to be shared

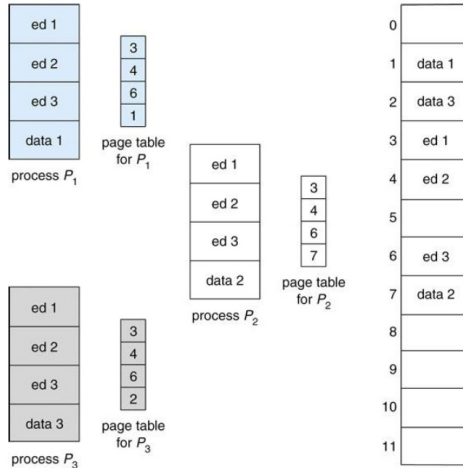
# Shared Data Segments [SGG12]

- Shared data segments are often implemented with
  - temporary, anonymous memory-mapped files
  - shared pages (with allocated space on backing store)





# Shared Pages Example [SGG12]

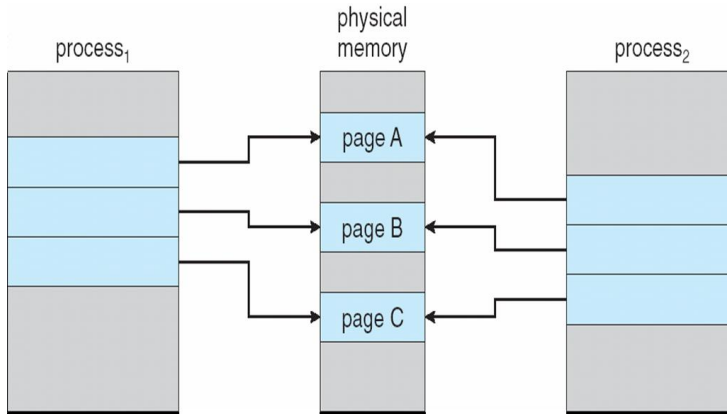


# Copy-On-Write

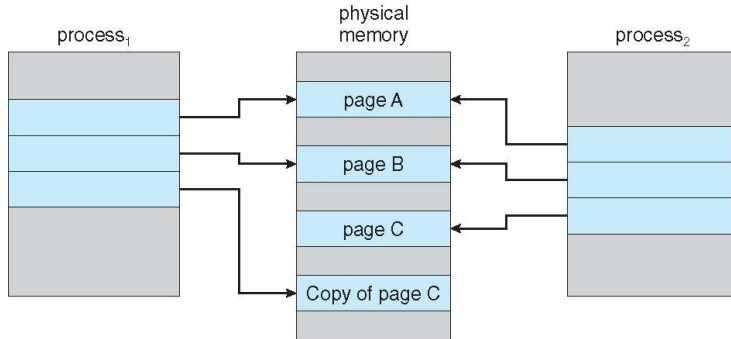
Copy-on-Write (COW) allows multiple processes to initially share the same pages in memory

- A page is copied only if one of the processes attempts to modify it
- COW allows more efficient process creation as only modified pages are copied

# COW: Before Process 1 Modifies Page C [SGG12]



# COW: After Process 1 Modifies Page C [SGG12]



# Frame Allocation

# Local vs. Global Allocation

- **Global allocation:** All frames are considered for replacement
  - Does not consider page ownership
  - One process can get another process's frame
  - Does not protect process from a process that hogs all memory
- **Local allocation:** Only frames of the faulting process are considered for replacement
  - Isolates processes (or users)
  - Separately determine how many frames each process gets

# Fixed Allocation

- **Equal allocation:** All processes get the same amount of frames
  - e.g., there are 100 frames and 5 processes → each process gets 20 frames
- **Proportional allocation:** Allocate according to the size of the process

$s_i$  = size of process  $p_i$

$$S = \sum s_i$$

$m$  = total number of frames

$a_i$  is the allocation for  $p_i$ :  $a_i = \frac{s_i}{S} \times m$

Example:

$$m = 64$$

$$s_1 = 10$$

$$s_2 = 127$$

$$a_1 = \frac{10}{137} \times 64 \approx 5$$


$$a_2 = \frac{127}{137} \times 64 \approx 59$$

# Priority Allocation

- **Priority Allocation** (global replacement)
  - Proportional allocation scheme using priorities rather than size
- If process  $P_i$  generates a page fault
  - Select one of its frames for replacement or
  - Select a frame from a process with lower priority



# Locality in a Memory-Reference Pattern [SGG12]



# Memory Locality

- Background storage is much slower than memory
  - Paging extends memory size using background storage
  - Goal: Run near memory speed, not near background storage speed
- **Pareto principle** applies to working sets of processes
  - 10% of memory gets 90% of the references
  - Goal: Keep those 10% in memory, the rest on disk
- Problem: How do we identify those 10%?


# Working-Set Model

- $\Delta$ : Working-set window
  - A fixed number of page references
  - Example: 10,000 instructions (#instruction = #page ref?)
- $WSS_i$ : Working set of process  $P_i$ 
  - Total number of pages referenced in the most recent  $\Delta$  (varies in time)
  - $\Delta$  too small: Will not encompass entire locality
  - $\Delta$  too large: Will encompass several localities
  - $\Delta = \infty$ : Will encompass entire program
- $D = \sum WSS_i$ : Total demand for frames
- If  $D > m$ : Thrashing
  - Policy: If  $D > m$ , suspend a process


# Working-Set Model [SGG12]

page reference table

... 2 6 1 5 7 7 7 7 5 1 6 2 3 4 1 2 3 4 4 4 4 3 4 3 4 4 4 4 1 3 2 3 4 4 4 4 3 4 4 4 ...



$$WS(t_1) = \{1, 2, 5, 6, 7\}$$




$$WS(t_2) = \{3, 4\}$$

## Example: Keeping Track of the Working Set


- The MMU automatically sets the **reference bit** in the respective page table entry every time a page is referenced
- Set a timer to scan all page table entries for reference bits
  - e.g.,  $\Delta = 10,000$
  - Timer interrupts after every 5,000 time units
  - Keep 2-bit history for each page in addition to the reference-bit
    - On timer interrupt, do for each page:
      - Shift reference bit into the 2-bit history
      - Reset reference bit
  - If history  $\neq 0$ : Page is in working set
  - Not accurate, because window is moving in large steps
    - Improvement: 10 bits and interrupt every 1000 time units

# Working Set and Page Fault Rates [SGG12]



# Page Fault Frequency Allocation Scheme

- Establish “acceptable” page fault rate
  - If actual rate too low, give frames to other process
  - If actual rate too high, allocate more frames to process [SGG12]



# Page Fetch Policy: Demand-Paging

- When should the OS allocate new pages?
  - Two possibilities: Pre-paging and demand-paging
- **Demand-Paging**: Transfer only pages that raise page faults
- + Only transfer what is needed
- + Less memory needed per process  
(higher degree of multiprogramming possible)
- “Many” initial page faults when a task starts
- More I/O operations → More I/O overhead



# Page Fetch Policy: Pre-Paging

- **Pre-Paging**: Speculatively transfer pages to RAM
  - At every page fault: speculate what else should be loaded
  - E.g., load entire text section when starting process
- + Improves disk I/O throughput by reading chunks
- Wastes I/O bandwidth if page is never used
- Can destroy the working set of other processes in case of page stealing

# References I

- [KELS62] T. Kilburn, D. B. G. Edwards, M. J. Lanigan, and F. H. Sumner. One-level storage system. *IRE Transactions on Electronic Computers*, EC-11(2):223–235, 1962.
- [SGG12] Abraham Silberschatz, Peter B. Galvin, and Greg Gagne. *Operating System Concepts*. Wiley Publishing, 9th edition, 2012.