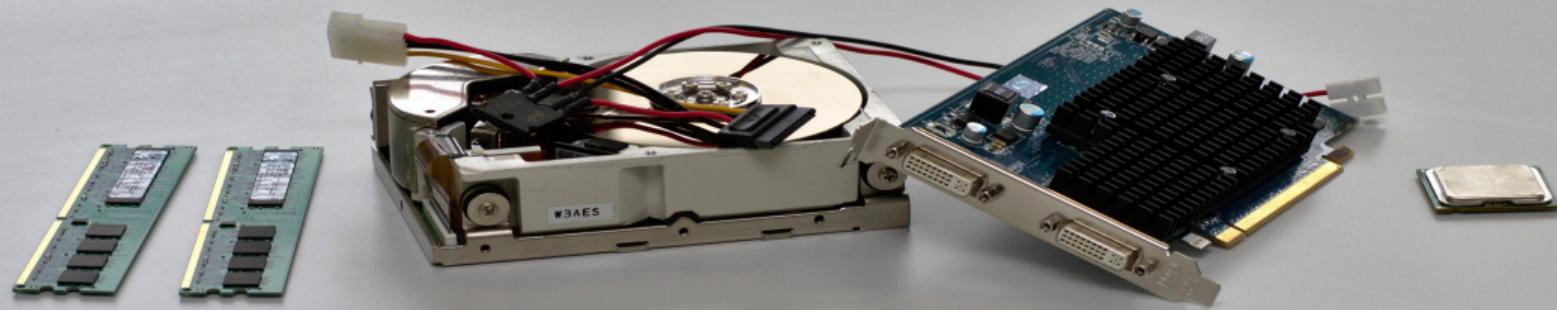


Operating Systems

15. Device Management

Prof. Dr. Frank Bellosa | WT 2020/2021

KARLSRUHE INSTITUTE OF TECHNOLOGY (KIT) – ITEC – OPERATING SYSTEMS



Device Management Objectives

- Abstraction from details of physical devices
- Uniform Naming that does not depend on HW details
- Serialization of I/O-operations by concurrent applications
- Protection of standard-devices against unauthorized accesses
- Buffering, if data from/to a device cannot be stored in the final destination
- Error Handling of sporadic device errors
- Virtualizing physical devices via memory and time multiplexing (e.g. pty, RAM disk)

Device Management Objectives

- Abstraction from details of physical devices
- Uniform Naming that does not depend on HW details
- Serialization of I/O-operations by concurrent applications
- Protection of standard-devices against unauthorized accesses
- Buffering, if data from/to a device cannot be stored in the final destination
- Error Handling of sporadic device errors
- Virtualizing physical devices via memory and time multiplexing (e.g. pty, RAM disk)

Device Management Objectives

- Abstraction from details of physical devices
- Uniform Naming that does not depend on HW details
- Serialization of I/O-operations by concurrent applications
- Protection of standard-devices against unauthorized accesses
- Buffering, if data from/to a device cannot be stored in the final destination
- Error Handling of sporadic device errors
- Virtualizing physical devices via memory and time multiplexing (e.g. pty, RAM disk)

Device Management Objectives

- Abstraction from details of physical devices
- Uniform Naming that does not depend on HW details
- Serialization of I/O-operations by concurrent applications
- Protection of standard-devices against unauthorized accesses
- Buffering, if data from/to a device cannot be stored in the final destination
- Error Handling of sporadic device errors
- Virtualizing physical devices via memory and time multiplexing (e.g. pty, RAM disk)

Device Management Objectives

- Abstraction from details of physical devices
- Uniform Naming that does not depend on HW details
- Serialization of I/O-operations by concurrent applications
- Protection of standard-devices against unauthorized accesses
- Buffering, if data from/to a device cannot be stored in the final destination
- Error Handling of sporadic device errors
- Virtualizing physical devices via memory and time multiplexing (e.g. pty, RAM disk)

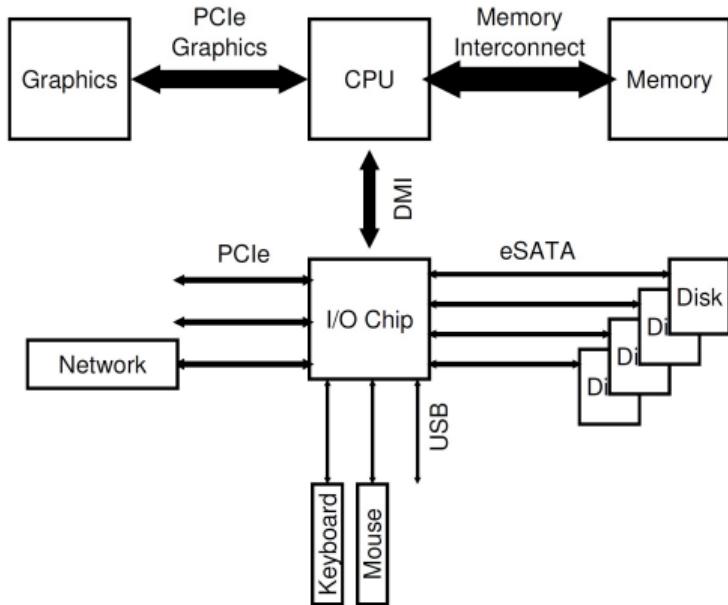
Device Management Objectives

- Abstraction from details of physical devices
- Uniform Naming that does not depend on HW details
- Serialization of I/O-operations by concurrent applications
- Protection of standard-devices against unauthorized accesses
- Buffering, if data from/to a device cannot be stored in the final destination
- Error Handling of sporadic device errors
- Virtualizing physical devices via memory and time multiplexing (e.g. pty, RAM disk)

Device Management Objectives

- Abstraction from details of physical devices
- Uniform Naming that does not depend on HW details
- Serialization of I/O-operations by concurrent applications
- Protection of standard-devices against unauthorized accesses
- Buffering, if data from/to a device cannot be stored in the final destination
- Error Handling of sporadic device errors
- Virtualizing physical devices via memory and time multiplexing (e.g. pty, RAM disk)

A Modern System Architecture



[ADAD18]

Characteristics of I/O Devices (1)

- **Block devices** include disk drives
 - Commands include read, write, seek
 - Raw I/O or file-system access
 - Memory-mapped file access possible
- **Character devices** include keyboards, mice, serial ports
 - Commands include get, put
 - Libraries layered on top allow line editing
- **Network devices** Specific communication patterns and addressing schemes require own interface
 - Usually addressed with socket interface
 - Separates network protocol from network operation

Characteristics of I/O Devices (1)

- **Block devices** include disk drives
 - Commands include read, write, seek
 - Raw I/O or file-system access
 - Memory-mapped file access possible
- **Character devices** include keyboards, mice, serial ports
 - Commands include get, put
 - Libraries layered on top allow line editing
- **Network devices** Specific communication patterns and addressing schemes require own interface
 - Usually addressed with socket interface
 - Separates network protocol from network operation

Characteristics of I/O Devices (1)

- **Block devices** include disk drives
 - Commands include read, write, seek
 - Raw I/O or file-system access
 - Memory-mapped file access possible
- **Character devices** include keyboards, mice, serial ports
 - Commands include get, put
 - Libraries layered on top allow line editing
- **Network devices** Specific communication patterns and addressing schemes require own interface
 - Usually addressed with socket interface
 - Separates network protocol from network operation

Characteristics of I/O Devices (2)

aspect	variation	example
data-transfer mode	character block	terminal disk
access method	sequential random	printer optical drive
transfer schedule	synchronous asynchronous	3D printer keyboard
sharing	dedicated sharable	tape network interface
device speed	latency / seek time transfer rate	
I/O direction	read only write only read-write	video capture printer disk

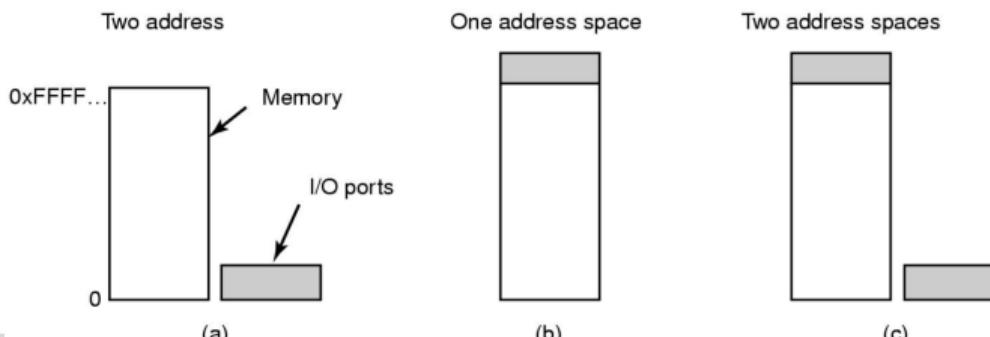
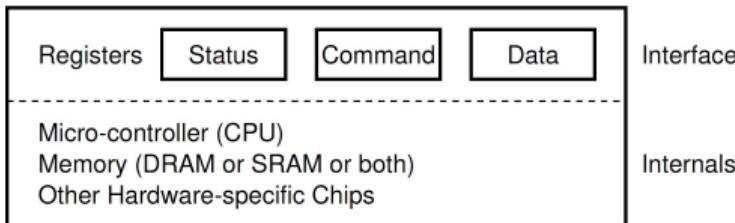
I/O Hardware

■ Canonical device interface [ADAD18]

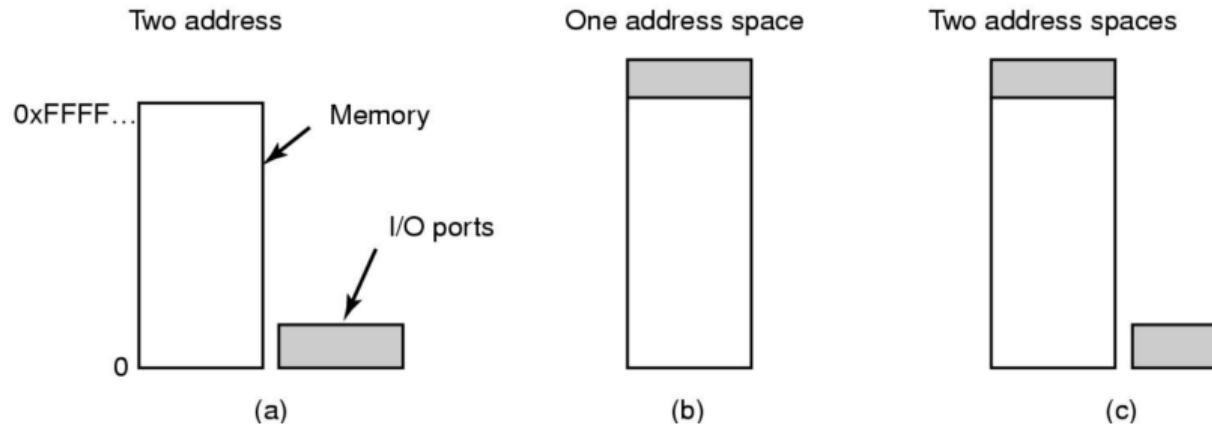
- Status register
- Control register
- Data-in register
- Data-out register

■ Devices are referenced by addresses [TB15]

- Direct I/O instructions (e.g. to access x86 I/O ports)
- Memory-mapped I/O



Ports v. Memory-Mapped I/O



- Separate I/O-address space and memory address space

- IN R0, 4 //<port 4> → R0
 - MOV R0, 4 //<4> → R0

- Memory-mapped I/O

- Hybrid (Intel)

Device I/O Port Locations on PCs (partial)

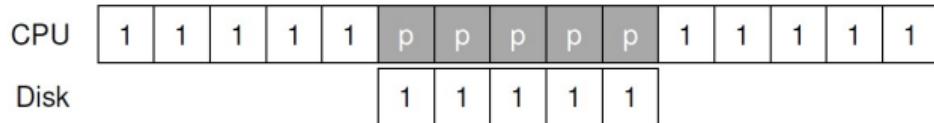
I/O address range (hexadecimal)	device
000 – 00F	DMA controller
020 – 021	interrupt controller
040 – 043	timer
200 – 20F	game controller
2F8 – 2FF	serial port (secondary)
320 – 32F	hard-disk controller
378 – 37F	parallel port
3D0 – 3DF	graphics controller
3F0 – 3F7	diskette-drive controller
3F8 – 3FF	serial port (primary)

[SGG12]

Polling vs. Interrupt-driven I/O [ADAD18]

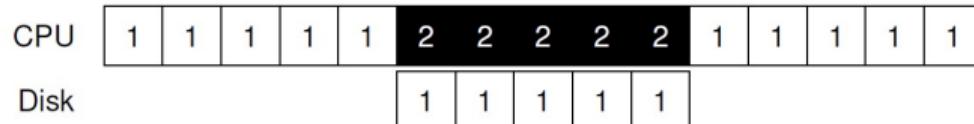
■ Programmed I/O

- Thread is busy-waiting for the I/O-operation
- Kernel thread is **Polling** the state of an I/O device
 - device ready
 - busy
 - error

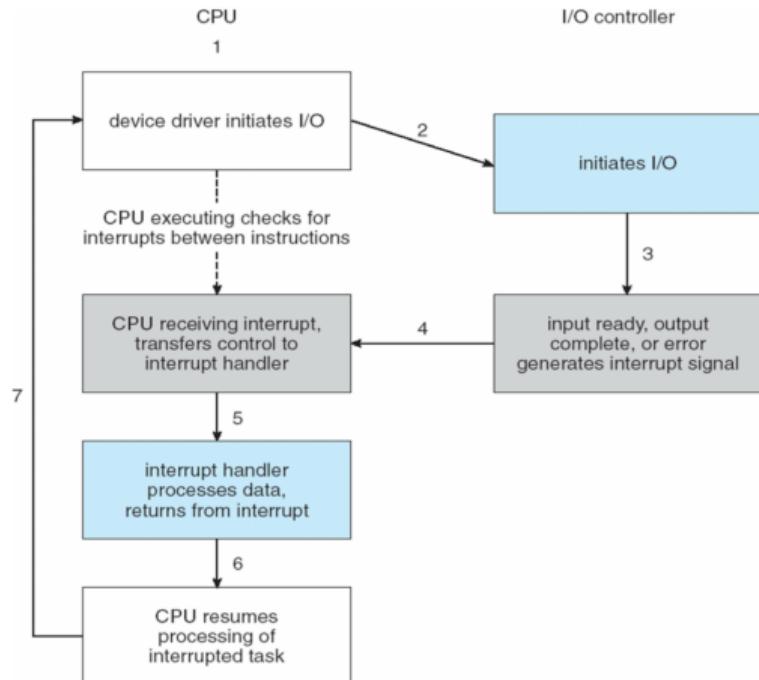


■ Interrupt-driven I/O

- I/O-command is issued
- processor continues executing instructions
- I/O-device sends an interrupt when I/O-command is done



Interrupt-Driven I/O Cycle



[SGG12]

Steps for Handling an Interrupt (1)

- ① Save registers not already saved by HW-interrupt mechanism
- ② Set up context (address space) for interrupt service procedure
 - Typically, handler runs in the context of the currently running process/task
⇒ no expensive context switch
- ③ Set up stack for interrupt service procedure
 - Handler usually runs on the kernel stack of the current process/kernel-level thread
 - Handler cannot block, otherwise the unlucky interrupted process/kernel-thread would also be blocked, might lead to starvation or even to a deadlock
- ④ Acknowledge/mask interrupt controller, thus re-enable other interrupts

Steps for Handling an Interrupt (2)

- 5 Run interrupt service procedure
 - Acknowledge interrupt at device level
 - Figures out what caused the interrupt, e.g.
 - Received a network packet
 - Disk read has properly finished, ...
 - If needed, it signals the blocked device driver
- 6 In some cases, we have to wake up a higher priority process/kernel level thread
 - Potentially schedule another process/kernel-level thread
 - Set up MMU context for process to run next
- 7 Load new/original process' registers
- 8 Return from interrupt, start running new/original process

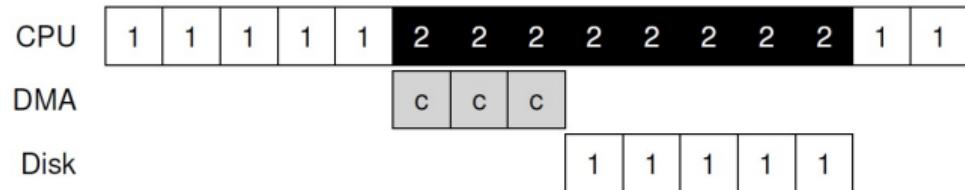
Polling vs. Interrupts [ADAD18]

- Fast device: better to spin than take interrupt overhead
- Device time unknown: Hybrid approach (spin then use interrupts)
- Flood of interrupts arrive
 - Can lead to livelock (always handling interrupts)
 - Better to ignore interrupts while make some progress handling them
- Interrupt coalescing (batch together several interrupts)

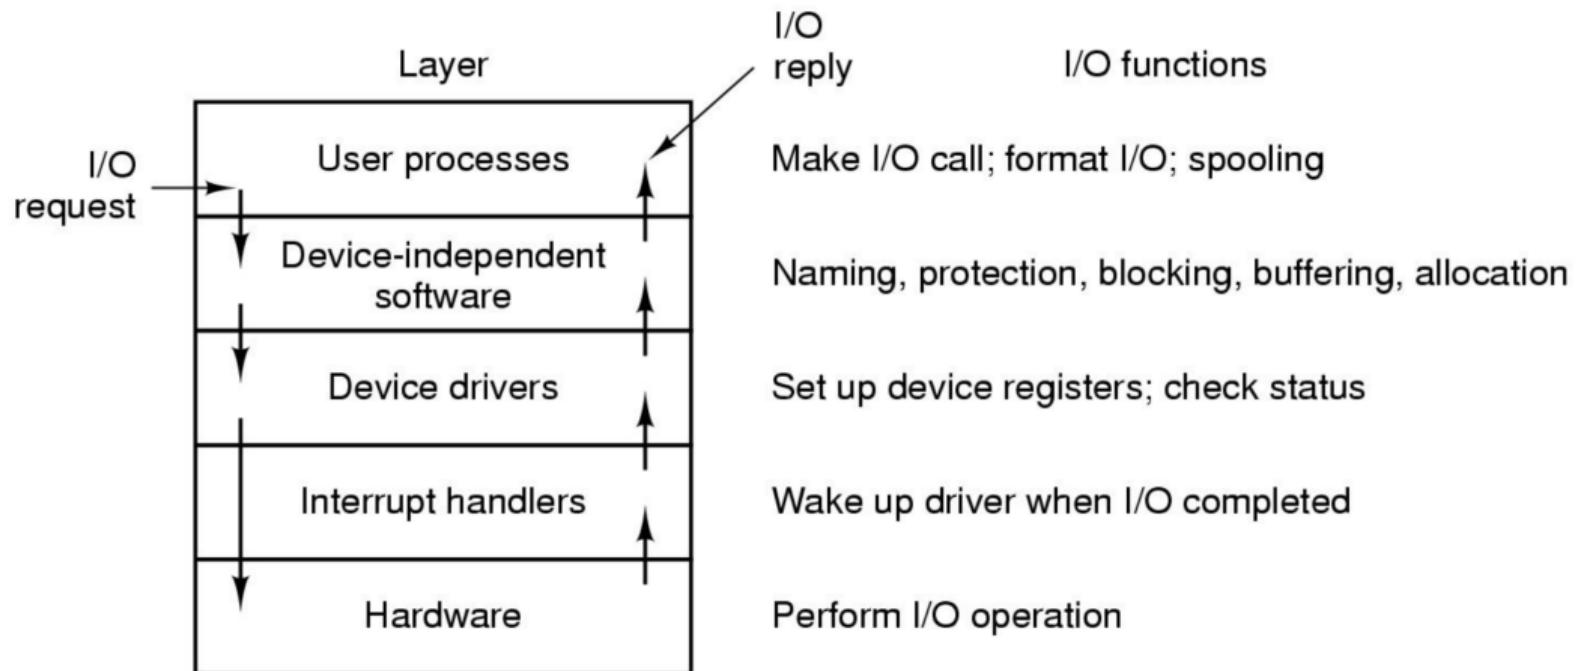
Direct Memory Access (DMA)

■ Direct Memory Access (DMA)

- DMA module controls exchange of data between main memory and I/O device
- processor interrupted after entire block has been transferred
- bypasses CPU to transfer data directly between I/O device and memory



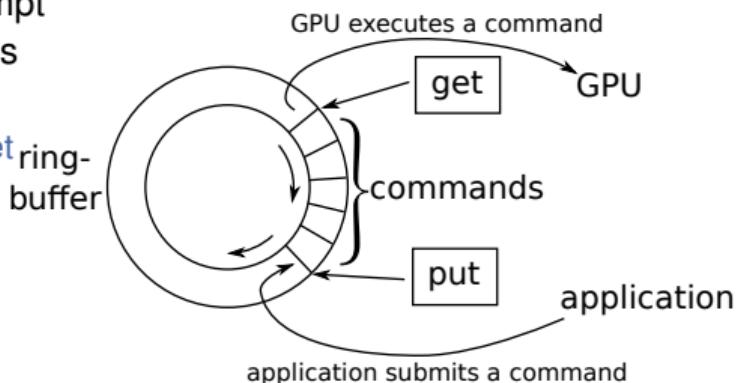
Layered I/O Subsystem



[TB15]

Functions of the Kernel I/O Subsystem (1)

- Device reservation – provides exclusive access to a device
 - System calls for allocation and deallocation
- Protection
 - I/O must be performed via system calls
 - Memory-mapped and I/O port memory locations must be protected
 - User process may accidentally or purposefully attempt to disrupt normal operation via illegal I/O instructions
 - Example: Nvidia General Purpose GPU
 - Command submission channel: **Ring-buffer** and **put/get** ring-buffer device registers are memory-mapped
 - Mapping can be exposed to application
 - Application can submit commands in user-mode



Functions of the Kernel I/O Subsystem (2)

- Spooling
 - Hold output for a device, if device can serve only one request at a time (i.e. printing)
- Uniform kernel interface for device code
 - Drivers use a defined interface to kernel service, e.g. kmalloc, install IRQ handler, etc.
 - Allows kernel to evolve without breaking device drivers

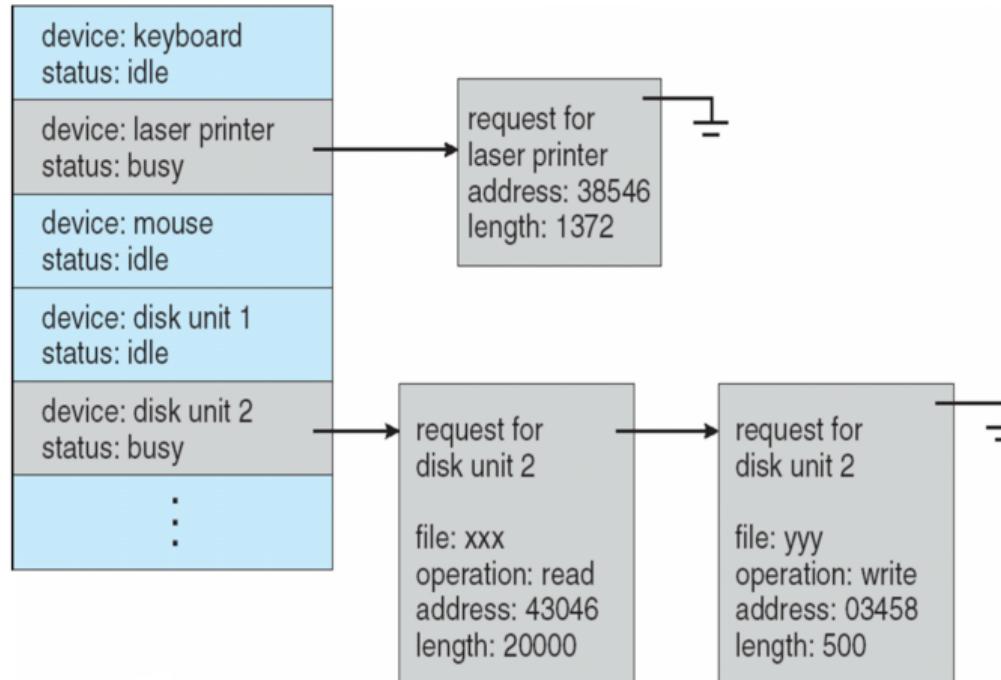
Functions of the Kernel I/O Subsystem (3)

- Request Scheduling
 - Some I/O request ordering via per-device queue
 - Some OSs try fairness
- Buffering – store data in memory while transferring between devices
 - To cope with device speed mismatch
 - To cope with device transfer size mismatch
 - To maintain “copy semantics”
- Error handling and reporting
 - OS can recover from errors e.g., disk read, device unavailable, transient write failures
 - Error reporting to upper levels, i.e. all errors the driver cannot resolve

Device Driver

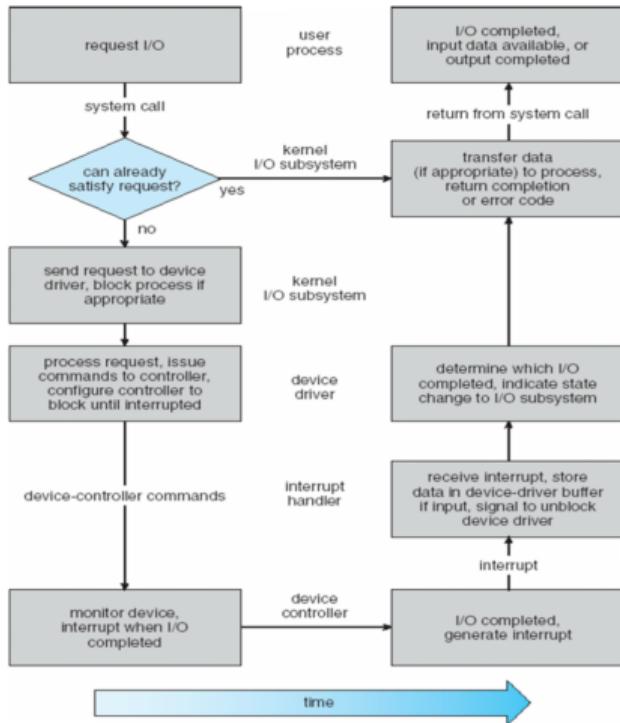
- After issue the command to the device, device either
 - completes immediately and the driver simply returns to the caller **or it**
 - processes request and the driver usually blocks waiting for an I/O (complete) interrupt signal
- Drivers are reentrant as they can be called by another process while a process is already blocked in the driver
 - Reentrant: code that can be executed by more than one thread (or CPU) at the same time
 - Manages concurrency using sync primitives

Device-status Table



[SGG12]

Life Cycle of an I/O Request



References I

- [ADAD18] R.H. Arpaci-Dusseau and A.C. Arpaci-Dusseau. *Operating Systems: Three Easy Pieces*. 2018.
retrieved from <http://pages.cs.wisc.edu/~remzi/OSTEP/>.
- [SGG12] Abraham Silberschatz, Peter B. Galvin, and Greg Gagne. *Operating System Concepts*. Wiley Publishing, 9th edition, 2012.
- [TB15] Andrew S Tanenbaum and Herbert Bos. *Modern operating systems*. Pearson, 4th edition, 2015.