

# Operating Systems

## 05. Segmentation

Prof. Dr. Frank Bellosa | WT 2020/2021

KARLSRUHE INSTITUTE OF TECHNOLOGY (KIT) – ITEC – OPERATING SYSTEMS



# Process Address Space Layout

**OS** Addresses of the kernel

**Stack** Local variables, function call parameters, return addresses

**Heap** Dynamically allocated data


**BSS** Uninitialized data

**Data** Initialized data

**RO-Data** Read-only data

**Text** Executable code

How does the OS organize and provision memory for multiple processes?



# Main Memory

- Main memory and registers are the only storage that the CPU can access directly
- Program must be brought into memory from background storage and placed within a process' address space for it to be run
- Early computers had no memory abstraction
  - Programs accessed physical memory directly
- Multiple processes can be run concurrently even without memory abstraction
  - Swapping
  - Static Relocation

# Swapping

- **Swapping** denotes

- saving a program's state on background storage (**roll-out**)
- replacing it with another program's state (**roll-in**) [SGG12]


- + Only needs hardware support to protect the kernel, but not to protect processes from one another

- Very slow

- Major part of swap time is transfer time
- Total transfer time is directly proportional to the amount of memory swapped


- At every point in time only one process runs: no parallelism

- This process owns the entire physical user space




# Overlays

- What if the process you want to run needs more memory than available?
- Need to partition program manually



# Static Relocation

- Another possibility to solve address conflicts when loading multiple processes is **static relocation**
  - The OS adds a fixed offset to every address in a program when loading it and creating a process from it
- Same address space (physical addresses) for every process
  - No protection: Every program sees and can access every address.
  - What if gcc needs more memory for its abstract syntax tree?
  - What if mplayer is pausing playback and currently doesn't need memory?  
Can it be reused by other processes?
  - What if no contiguous free region fits program?



# Desired properties when sharing physical memory

## ■ Protection

- A bug in one process must not corrupt memory in another
- Don't allow processes to peek into other processes' memory

## ■ Transparency

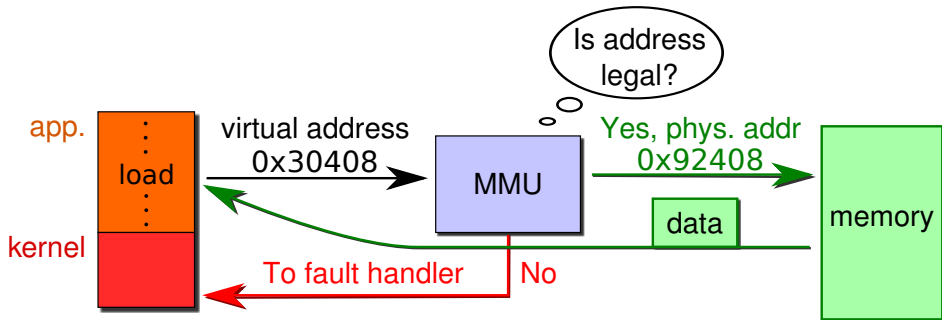
- A process shouldn't require particular physical memory addresses
- Processes should be able to use large amounts of contiguous memory

## ■ Resource exhaustion

- Allow that the sum of sizes of all processes is greater than physical memory

# Ideal Memory-Management Unit (MMU)

- Architectural support to achieve safe and secure protection
- Hardware device maps virtual to physical address
- The user program deals with virtual addresses
  - It never sees the real physical addresses

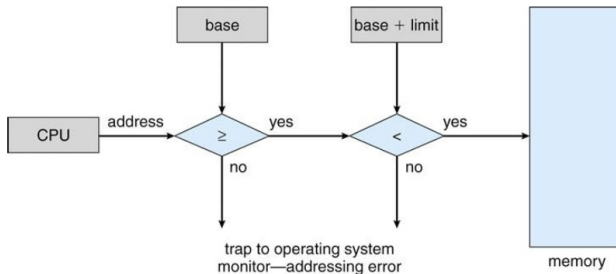


[Maz20]




# Base and Limit Registers

- Provides protection with static relocation
  - Introduce special **base** and **limit** registers
- On every load/store the MMU
  - Checks if the address is larger or equal to **base**
  - Checks if the address is smaller than **base + limit**
  - Use the CPU address as the physical address in memory [SGG12]



# Protecting the Kernel with Base and Limit Registers

- Need to protect OS from processes
- Main memory split into two partitions
  - Resident operating system, usually held in **low memory**
  - User processes held in **high memory**
- OS can access all process partitions e.g., to copy syscall parameters
- MMU denies processes access to OS memory [SGG12]




# Base and Limit Registers

- + Straight forward to implement MMU
  - Only need to load new base and limit registers to switch address space
- + Very quick at run-time
  - Only two comparisons (can do both in parallel)
  - Compute **base** + **limit** in advance
- How do you grow a process' address space?
- How do you share code or data?

# Segmentation

- Solution for shortcomings of Base + Limit approach:
  - Use **multiple** Base+Limit register pairs **per process**

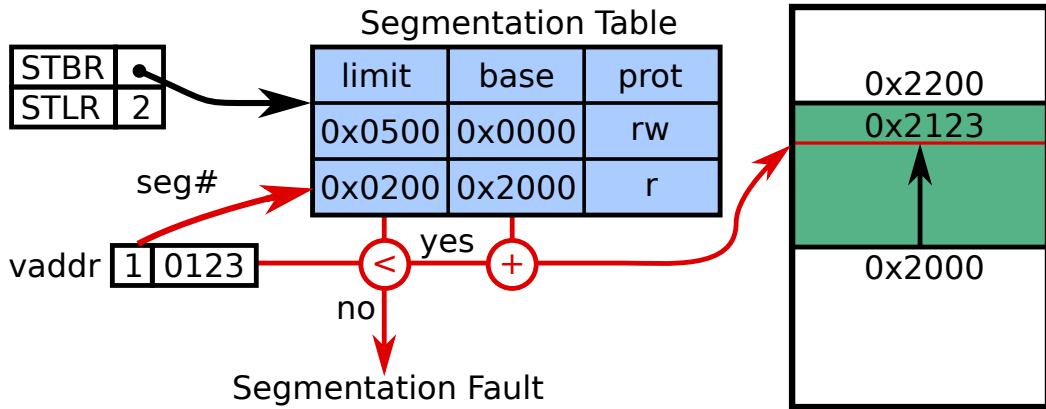


→ Can keep some **segments** private, share others

# Segmentation Architecture

- Virtual address consists of a tuple:  $\langle \text{segment \#}, \text{offset} \rangle$ 
  - Can be encoded in the address (PDP-10 – seg #: high bits, offset: low bits)
  - Can be selected by an instruction or an operand
- Each process (address space) has a **segment table** that maps virtual address to physical addresses in memory
  - **Base** Starting physical address where the segment resides in memory
  - **Limit** Length of the segment
  - **Protection** Access restriction (read/write) to make safe sharing possible
- The MMU has two registers that identify the current address space
  - **Segment-table base register (STBR)** points to the segment table location of the current process
  - **Segment-table length register (STLR)** indicates number of segments used by the process (segment # is legal if it is  $< \text{STLR}$ )

# Segmentation Mechanics




[Maz20]

# Segmentation Trade-offs

- + Dynamic relocation is possible  
(adaption of base in segment-table)
- + Makes data/code sharing between processes possible without compromising confidentiality and safety/security
- + Process doesn't need one large contiguous physical memory area  
→ easier placement
- + Don't need entire process in memory  
→ Can overcommit memory with segment swap after a segmentation fault
- Segments need to be kept contiguous in physical memory
- **Fragmentation** of physical memory

# External Fragmentation

- Fragmentation  $\equiv$  The inability to use free memory
- External Fragmentation Sum of free memory satisfies requested amount of memory. Request cannot be fulfilled with contiguous memory.
- Compaction reduces external fragmentation
  - Close gaps by moving allocated memory in one direction
  - Results in a large free block on the other side
  - Compaction is possible only if relocation is dynamic, and can be done at execution time.



- Expensive: Need to halt process while moving data and updating tables.



# References I

- [Maz20] David Mazières. Cs140 – operating systems. *Stanford University*, 2020.
- [SGG12] Abraham Silberschatz, Peter B. Galvin, and Greg Gagne. *Operating System Concepts*. Wiley Publishing, 9th edition, 2012.