

Operating Systems

06. Dynamic Memory Allocation

Prof. Dr. Frank Bellosa | WT 2020/2021

KARLSRUHE INSTITUTE OF TECHNOLOGY (KIT) – ITEC – OPERATING SYSTEMS



Dynamic Memory Allocation


- **Dynamic Memory Allocation:** Allocate and free memory chunks of “arbitrary size” at arbitrary points in time
 - Almost every program uses it (heap)
 - Complex data structures don’t have to be specified statically
 - Data can grow as a function of input size
 - Kernel itself uses dynamic memory allocation for its data structures and buffers
 - Implementation of dynamic memory allocation has a huge impact on performance
 - Both in user space and in kernel
 - Common knowledge: It is impossible to construct a memory allocator that always performs well
 - “For any possible allocation algorithm, there exists a stream of allocation and deallocation requests that defeat the allocator and force it into severe fragmentation” [Rob71]
- Picking a good allocation strategy requires a trade-off

What does a Dynamic Memory Allocator do?

- Initially has a pool of free memory
- Needs to satisfy arbitrary `allocate` and `free` requests from that pool
- Cannot control the order or the number of requests
- Cannot move allocated regions (no compaction!)
 - Relocation is not possible (e.g., within a virtual AS)
 - Bad allocation decision is permanent!
- Fragmentation is a core problem
- Needs to track which parts are in use and which parts are free


Bitmap

- Divide memory in allocation units of fixed size
- Use a bitmap to keep track if allocated (1) or free (0)
- Needs additional data structure to store allocation length
 - Otherwise cannot infer whether two adjacent allocations belong together or not from bitmap, but need this info for **free**. (e.g., try **freeing** address 8) [TB15]



List

- Either: Use one list-node for each allocated area
 - Needs extra space for the list
 - Allocation lengths already stored
- Or: Use one list-node for each unallocated area
 - Can keep list in the unallocated area (e.g., store size of free area and pointer to next free area in the free area itself)
 - Can search for free space with low overhead
- Or: Both (Double linked list speeds up merging after freeing) [TB15]




Why is Dynamic Memory Allocation Hard?

- Fragmentation is hard to handle
 - **Fragmentation**: Inability to use free memory
 - **External fragmentation**: Sum of free space is sufficient but cannot allocate sufficiently large contiguous block of free memory
 - **Internal fragmentation**: Overallocate resource requests to align memory blocks. Don't have free blocks left although there is sufficient unused memory within the blocks.
- Three factors required for fragmentation to occur
 - Different lifetimes (symmetric allocation times: no problem → stack)
 - Different sizes (same size: no problem → next allocation fits into any hole)
 - Inability to relocate previous allocations

A Pathological Allocation Example

- All memory is allocated in 32-byte chunks
- Every other allocation is released
- If it now wants to allocate 33-bytes:
 - The resource request fails...
 - ...although half of the memory is free



- Required “gross” memory in bad allocator: $M \cdot \frac{n_{\max}}{n_{\min}}$
 - M = bytes of live data
 - n_{\min} = smallest allocation, n_{\max} = largest allocation
use maximum size for any allocation

Best Fit vs. Worst Fit

- Idea: Keep large free memory chunks together for larger allocation requests that may arrive later
- **Best-fit**: Allocate the smallest free block that is large enough to fulfill the allocation request
 - Must search entire list, unless ordered by size
 - During free: coalesce adjacent blocks
- Problem: **Sawdust**
 - External fragments of allocations are so small that over time leftovers with unusable sawdust are everywhere
- Idea: Minimize sawdust by turning the strategy around
- **Worst-fit**: Allocate the largest free block
 - Must also search entire list, unless ordered by size
- In reality: Worse fragmentation than best-fit

First Fit

- Idea: If you produce fragmentation with best fit and worst fit alike, try to optimize for allocation speed
- **First-fit**: Allocate the first hole that is big enough
 - Fastest allocation policy
 - Produces leftovers of variable size
- Pathological case: Mix short lived $2n$ -byte allocations with long-lived $(n+1)$ -byte allocations
 - Each time a large object is freed, a small chunk will be quickly taken, leaving a useless fragment
- In reality: Almost as good as best-fit

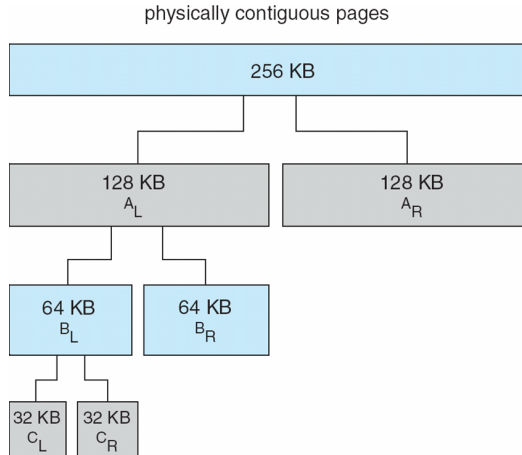
First Fit Nuances

- **First-fit sorted by address order:**
 - Blocks at front preferentially split, one at back only split when no larger one found before them
 - Seems to roughly sort free list by size
 - Similar to best fit
 - Sorting of list forces a large request to skip over many small blocks
- **LIFO First-fit:** Put object on front of list
 - Cheap & fast allocation policy
 - Hope same size used again (good cache locality)
- **Next fit:** Use First-fit, but remember where we found the last thing and start searching from there
 - Tends to break down entire list
 - Bad cache locality

Buddy Allocator

- Can be used to dynamically allocate contiguous chunks of fixed-size segments
 - e.g., used in the Linux kernel to allocate physical memory
- Allocates memory in powers of 2
 - All contiguous allocated/free memory chunks have fixed power-of-2 size
 - Request rounded up to next-higher power of 2
 - All chunks are **naturally aligned**
(i.e., their starting address is a multiple of their size)
- If no sufficiently small memory block is available
 - Select larger available chunk and split it into two equal-sized “buddies”
 - Continue until appropriately sized chunk is available
- If two buddies are both free
 - Merge buddies to larger chunk encompassing both buddies

Buddy System (2)




[SGG12]

Known patterns of real programs [Maz20]

- Most real programs exhibit 1 or 2 (or all 3) of the following patterns of alloc/dealloc:

- *Ramps*: accumulate data monotonically over time




- *Peaks*: allocate many objects, use briefly, then free all



- *Plateaus*: allocate many objects, use for a long time




Pattern 1: Ramps [Maz20]



- In a practical sense: ramp = no free!
 - Implication for fragmentation?
 - What happens if you evaluate allocator with ramp programs only?


Pattern 2: Peaks [Maz20]




- Peaks: allocate many objects, use briefly, then free all
 - Fragmentation a real danger
 - What happens if peak allocated from contiguous memory?
 - Interleave peak & ramp? Interleave two different peaks?

Exploiting Peaks [Maz20]

- Peak phases: alloc a lot, then free everything
 - So have new allocation interface: alloc as before, but only support free of everything
 - Called “arena allocation”, “obstack” (object stack), or `alloca` procedure call
- Arena = a linked list of large chunks of memory
 - Advantages: alloc is a pointer increment, free is “free”
No wasted space for tags or list pointers



Pattern 3: Plateau [Maz20]




trace of perl running a string processing script

- Peaks: allocate many objects for a long time
 - What happens if overlap with peak or different plateau ?

Known patterns of real programs [Maz20]


- Segregation = reduced fragmentation:
 - Allocated at same time ~ freed at same time
 - Different type ~ freed at different time



- Implementation observations :
 - Programs allocate small number of different sizes
 - Fragmentation at peak use more important than at low
 - Most allocations small (< 10 words)

SLAB Allocator

- Kernel often allocates/frees memory for few, specific data objects of fixed size
- A **slab** is made up of multiple pages of contiguous physical memory
- A **cache** consists of one or multiple slabs
- Each cache stores only one kind of object (fixed size)
- Linux uses Buddy Allocator as underlying allocator for slabs



[SGG12]

References I

- [Maz20] David Mazières. Cs140 – operating systems. *Stanford University*, 2020.
- [Rob71] J. M. Robson. An estimate of the store size necessary for dynamic storage allocation. *J. ACM*, 18(3):416–423, July 1971.
- [SGG12] Abraham Silberschatz, Peter B. Galvin, and Greg Gagne. *Operating System Concepts*. Wiley Publishing, 9th edition, 2012.
- [TB15] Andrew S Tanenbaum and Herbert Bos. *Modern operating systems*. Pearson, 4th edition, 2015.