


Operating Systems

09. Page Replacement


Prof. Dr. Frank Bellosa | WT 2020/2021

KARLSRUHE INSTITUTE OF TECHNOLOGY (KIT) – ITEC – OPERATING SYSTEMS



Naïve Page Replacement [SGG12]

- Save/clear victim page
 - Drop page if fetched from disk (e.g., code) and clean (PTE dirty bit)
 - Write back modifications if from disk and dirty
 - Write pagefile/swap partition otherwise (e.g., stack, heap memory)
- Unmap page from old AS
 - unset valid bit in PTE
 - flush TLB
- Prepare the new page
 - e.g., NULL page
 - e.g., load new contents
- Map the page frame into the new address space(s)
 - set valid bit in PTE
 - flush TLB



Page Buffering


- Problem: Naïve page replacement encompasses two I/O transfers swapping out (**demand cleaning**) and swapping the new page in
 - Both operations block the page fault from completing
- Goal: Reduce I/O from critical page fault path to speed up page faults
- Idea: Keep pool of free page frames (**pre-cleaning**)
 - On a page fault, use a page frame from the free pool
 - Run a daemon that cleans (write back changes), reclaims (unmap), and scrubs (zero out) pages for the free pool in the background
- Such a free pool smoothes out I/O and speeds up paging significantly
- Remaining problem: Which pages to select as victims?
 - Goal: Identify a page that has left the working set of its process to add it to the free pool
 - Success metric: Low overall page fault rate

First-In-First-Out (FIFO) Page Replacement

- Evict the oldest fetched page in the system

reference string


7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1



page frames

Intuition: Page Fault Rate vs. Number of Frames

- Intuitively one would say that the page fault rate decreases when the amount of memory increases
- This is true most of the time, but not universally



Belady's Anomaly [Bel66]

- Reference string for all our examples: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5
- Belady's Anomaly
 - When using FIFO page replacement, for every number N of page frames you can construct a reference string that performs worse with $N+1$ frames
 - With FIFO it is possible to get more page faults with more page frames

3 frames

1	1	4	5
2	2	1	3
3	3	2	4

9 page faults


4 frames

1	1	5	4
2	2	1	5
3	3	2	
4	4	3	

10 page faults


Belady's Anomaly using FIFO page replacement

- More physical memory doesn't always imply fewer faults [SGG12]



Oracle: Optimal Page Replacement (OPT)

- The optimal page replacement strategy is to replace the page whose next reference is furthest in the future [SGG12]



- Example with 4 frames:
1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5
 - Cannot predict the future
 - Not suitable in practice
 - However: Good metric to check how well other algorithms perform
- | |
|---|
| 2 |
| 3 |
| 4 |

5
- 6 page faults



Least Recently Used (LRU) Page Replacement

- Goal: Approximate Oracle page replacement
- Idea: Past often predicts the future well
- Assumption: Page used furthest in past is used furthest in the future

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2		2		4	4	4	0			1		1		1		
	0	0	0		0		0	0	3	3			3		0		0		
		1	1		3		3	2	2	2			2		2		7		

page frames

- Reference string
1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5
- 8 page faults

1	1	1	1	5
2	2	2	2	2
3	5	5	4	4
4	4	3	3	3

LRU: Easy to understand, hard to implement well

■ Cycle counter implementation

- Have MMU write CPU's time stamp counter to PTE on every access
- On a page fault: Scan all PTEs to find oldest counter value
- + Cheap at access if done in HW
- Memory traffic for scanning

■ Stack implementation

- Keep a doubly linked list of all page frames
- Move each referenced page to tail of list
- + Can find replacement victim in $O(1)$
- Need to change 6 pointers at every access

■ No silver bullet



- Observation: Predicting the future based on the past is not precise
- Conclusion: Relax requirements – maybe perfect LRU is not needed?
→ Approximate LRU

LRU Approximation: Clock Page Replacement

- Clock page replacement a.k.a. **second chance** page replacement
- Precondition: MMU sets **reference bit** in PTE at access ¹
 - Supported natively by most hardware (e.g., IA-32, x86-64, ...)
 - Can easily be emulated in systems with software managed TLB (e.g., MIPS)
- Keep all pages in circular FIFO list
- When searching for a victim scan pages in FIFO's order
 - If reference bit is 0 → use page as victim and advance hand
 - If reference bit is 1 → set to 0 and continue scanning
- Large memory → most pages referenced before scanned
 - Use 2 arms: Leading arm clears reference bit, trailing arm selects victim

¹Reference bit will be set for the first time by MMU after page fault when retrying access

Clock Page Replacement




Other replacement strategies

- Random eviction
 - Just pick a victim at random
 - Dirt simple and in reality not overly horrible
- Use larger counter: Use n-Bit reference counter instead of reference bit
 - Least frequently used (LFU)
 - Idea: Rarely used page is not in a working set
 - Replace page with smallest count
 - Problem: How to forget references of the past
 - Most frequently used (MFU)
 - Idea: The page with the smallest count was probably just brought in and will be used soon
 - Replace page with the largest count
 - Neither LFU nor MFU are used (no such hardware + far away from OPT)

Thrashing

- **Thrashing:** The system is busy swapping pages in and out
 - Each time one page is brought in, another page, whose contents will soon be referenced, is thrown out
 - Effect: Low CPU utilization, processes wait for pages to be fetched from disk
 - Consequence: OS “thinks” that it needs higher degree of multiprogramming [SGG12]




Reasons for Thrashing

- Memory too small to hold hot memory of a single process (the 10%)
- Access pattern has no temporal locality
 - Process doesn't follow the pareto principle
 - Past \neq future
- Each process fits individually, but too many for system
 - Degree of multiprogramming too high
- Page replacement policy doesn't work well

Solaris

- Maintains a list of free pages to assign to faulting processes
 - Paging is performed by *pageout* process
 - Scans pages using modified clock algorithm
 - *Scanrate* ranges from *slowscan* to *fastscan*
 - Free memory thresholds determine the behavior of *pageout*
 - *Lotsfree*: Begin paging
 - *Desfree*: Increase paging
 - *Minfree*: Begin swapping
- [SGG12]



References I

- [Bel66] L. A. Belady. A study of replacement algorithms for a virtual-storage computer. *IBM Systems Journal*, 5(2):78–101, 1966.
- [SGG12] Abraham Silberschatz, Peter B. Galvin, and Greg Gagne. *Operating System Concepts*. Wiley Publishing, 9th edition, 2012.