


Operating Systems

10. CPU Caches

Prof. Dr. Frank Bellosa | WT 2020/2021


KARLSRUHE INSTITUTE OF TECHNOLOGY (KIT) – ITEC – OPERATING SYSTEMS



Managing Memory

- Ideal memory properties
 - Large
 - Fast
 - Nonvolatile
 - Cheap
- Real memory: trade off properties above
 - Tape/Disk = large + slow + cheap + nonvolatile
 - SSD = not too large + not too cheap + nonvolatile + low endurance
 - PCRAM, Optane® = fast (Read) + slow (write) + nonvolatile
 - DRAM, SRAM = fast + volatile + expensive
 - Registers = very fast + volatile + very expensive
 - ...

CPU-Cache Location [Sch99]



- Buffer memory for exploiting temporal and spatial locality
- Low latency, high bandwidth
- Reduction of main memory accesses
- Reduction of memory bus traffic (important for multiprocessor systems)
- Buffer for asynchronous prefetch operations

Types of Cache Misses

■ Compulsory miss

- Cold start, first reference
- Data block was not cached before

■ Capacity miss

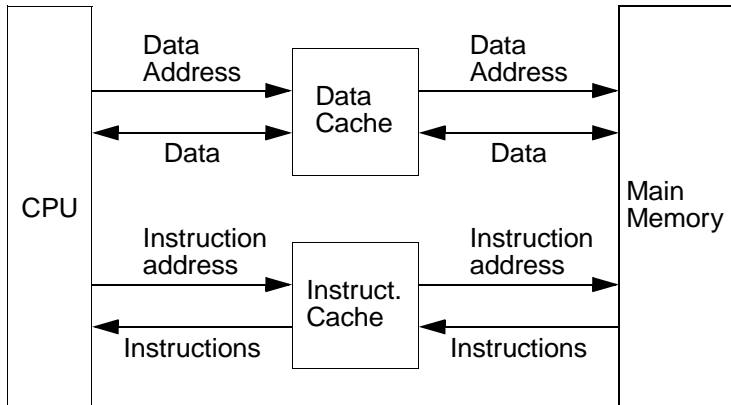
- Not all required data fits into the cache
- Accessed data was previously evicted to make room for different data

■ Conflict miss

- Collision, interference
- Depending on the cache organization, data items interfere with each other
- Fully associative caches are not prone to conflict misses

Harvard Architecture[Sch99]

- Separate buffer memory for data and instructions



Write and Replacement Policies

■ Cache hit

■ Write-through

- Main memory is always up-to-date
- Writes may be slow

■ Write-back

- Data is written only to the cache
- Main memory temporarily inconsistent

■ Cache miss

■ Write-allocate:


- To-be-written data item is read from main memory into the cache
- Write performed afterwards according to the write policy

■ Write-to-memory:

- Modification is performed only in main memory

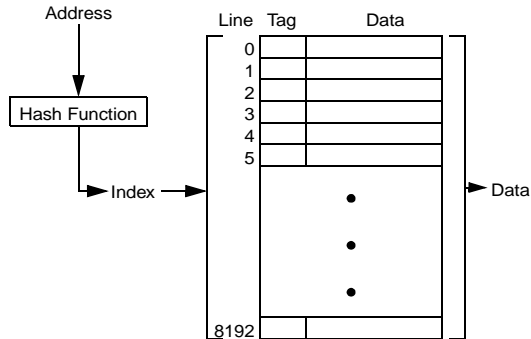
Cache Organization: Fully Associative Cache[Sch99]

- Cache-lines with fixed length (e.g., 32/64 Bytes)
- Data identified by tag-field
- Memory access: searching all tags for address tag




Cache Organization: Direct Mapped Cache[Sch99]

- Cache-lines with fixed length (e.g., 32/64 Bytes)
- Mapping from address to cache lines
- Data identified by tag-field



Cache Organization: Set Associative Cache[Sch99]




Example: 2-way set associative cache



- If $\#sets == \#lines \rightarrow$ Direct Mapped Cache
- If $\#sets == 1 \rightarrow$ Fully Associative Cache

Hash Functions[Sch99]

- **Modulo hashing**: e.g., bits 0 to 5 as byte-offset within a cache line, bits 6 to 18 to select a cache line, bits 19 to 31 as tag.
- **Arbitrary cut-out** of address to select cache line, e.g. bits 10 to 22. Suboptimal mapping of addresses to cache lines and therefore not used.





Cache Design Parameters


- Size and set size
 - Small cache → set-associative implementation with large sets
- Line length
 - Spatial locality → long cache lines
- Write policy
 - Temporal locality → write-back
- Replacement policy
- Using virtual or physical addresses for tagging/indexing

Long Cache Lines vs. Cache Alignment

- Cache line length influences what a good size for data structures is
- Problem: Multiple misses caused by “non-aligned” data structures spanning multiple cache lines
- Solution: Padding of structures to a multiple of line length

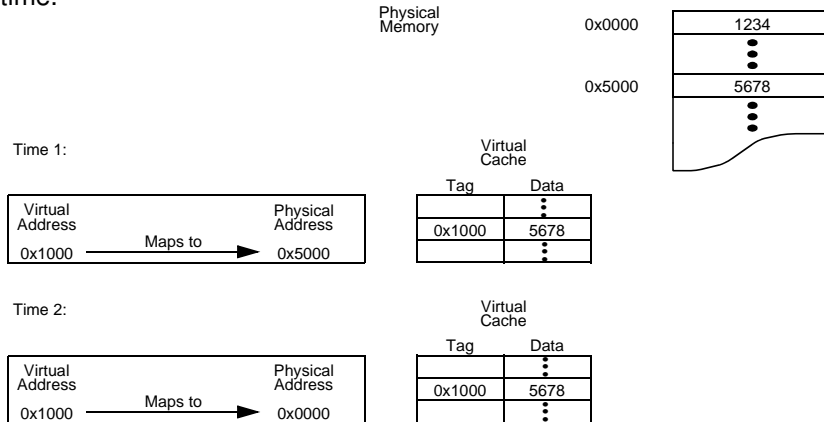


Virtually Indexed, Virtually Tagged (ARMv4/v5)[Sch99]




Ambiguity Problem[Sch99]

- Identical virtual addresses point to different physical addresses at different points in time.



Alias Problem[Sch99]

- Different virtual addresses point to the same physical memory location



After references to
0x2000 and 0x4000

Tag	Data
	⋮
0x2000	1234
	⋮
0x4000	1234
	⋮

Virtual
Cache

After modification
of 0x2000

Tag	Data
	⋮
0x2000	5678
	⋮
0x4000	1234
	⋮

Virtual
Cache

Virtually Indexed, Virtually Tagged Cache Operations


- **Context switch:** Cache must be invalidated (and written back if write-back is used), as identical virtual addresses of different address spaces might (likely) reference different physical addresses
- **fork:** The child needs a complete copy of the parent's address space. If copy operation runs in the context of the parent, special cache management is required.
- **exec:** Invalidate cache to prevent ambiguities. Not necessary to write content back, as memory is overwritten anyway.
- **exit:** Flush cache
- **brk/sbrk:** Growing requires no action, shrinking requires (selective) cache invalidations

Virtually Indexed, Virtually Tagged Cache Operations

- Shared memory and memory mapped files → alias problem (multiple virtual address refer to the same physical memory)
 - Disallow
 - Do not cache
 - Only allow addresses that map to the same cache line (if cache is direct-mapped and uses write-allocate)
 - Each frame accessible from exactly one virtual address at each point in time (requires invalidation of all alias pages in the page table)



Virtually Indexed, Virtually Tagged I/O [Sch99]

- Buffered I/O
 - No problems
- Unbuffered I/O
 - Write: Information might still be in the cache → write back before I/O starts
 - Read: Cache must be invalidated




Virtually Indexed, Virtually Tagged: Unbuffered I/O [Sch99]

- Additional problems if I/O region is not aligned with length of cache line




Virtually Indexed, Physically Tagged (ARMv6/v7)[Sch99]



- Often used as first-level caches
 - e.g., UltraSPARC II: 16 kB direct mapped
- Cache management in VIPT caches
 - No ambiguities
 - No cache flush on context switch
 - Shared memory/memory mapped files:
 - Virtual starting addresses must be mapped to the same cache line
 - I/O: Cache flush required as with virtually indexed, virtually tagged cache

VIPT Conflicts


- Data structures whose address distance is a multiple of the cache size are mapped to the same cache line



VIPT Runtime Properties

- Cache flush can be avoided most of the time
 - Fast context switches, interrupt-handling and system calls
- Deferred write-back after context switch
 - Avoids write accesses (performance gain)
 - Variable execution time caused by compulsory misses (depends on access pattern of previous thread)
- Variable execution time in case of dynamic memory management caused by conflict misses
- Variable search time caused by address translation in MMU
- Problematic in multiprocessor systems with shared memory:
Which line to invalidate?
 - Cache size is a small multiple of page size (factor 1-4)
 - Requires to only invalidate/flush 1-4 cache lines by cache coherency HW


Physically Indexed Physically Tagged Caches[Sch99]



- + Completely transparent to processor
- + No performance-critical system support required (including I/O)
- + SMPs with shared memory can use coherency protocol implemented in hardware

PIPT Random Allocation Conflicts

- Contiguous virtual memory is normally mapped to arbitrary free physical pages
- Page conflicts caused by random allocation of physical memory



PIPT Random Coloring Conflicts


- Consequences of random page coloring:
 - Cache Conflicts
 - Cache only partially used
 - Significant variations in runtime

PIPT Conflict Mitigation

- Sequential page colors for individual memory segments
 - E.g., red-yellow-green-blue-red-yellow . . .
- Cache partitioning [LHH97]
 - Divide physical memory in disjoint subsets.
 - All pages of a subset are mapped to the same cache partition
 - Example:
 - All red and blue pages for operating system
 - All yellow pages for “the” real-time application
 - All green pages for background processes

PIPT Conflict Mitigation: Intel® Cache Allocation Technology


- Class of service (CLOS) defines the amount of resources (cache space) available via MSRs(Model Specific Register).
- Associate each logical thread with an available logical CLOS
- As the OS swaps a thread onto a core, update the CLOS on the core via a CLOS-specific MSR



[Ngu16] [int20]

PIPT Conflict Mitigation: Page Recoloring

- Analysis of access patterns (WS analysis) and page-recoloring



References I

- [int20] *Four-Volume Set of Intel® 64 and IA-32 Architectures Software Developer's Manuals*, 2020. retrieved from <https://software.intel.com/content/www/us/en/develop/articles/intel-sdm.html#three-volume>.
- [LHH97] J. Liedtke, H. Härtig, and M. Hohmuth. Os-controlled cache predictability for real-time systems. In *Proceedings Third IEEE Real-Time Technology and Applications Symposium*, pages 213–224, 1997.
- [Ngu16] K. T. Nguyen. Usage models for cache allocation technology in the intel® xeon® processor e5 v4 family. *Intel Development Topics & Technologies*, 2016.
- [Sch99] Curt Schimmel. *UNIX systems for modern architectures : symmetric multiprocessing and caching for kernel programmers*. Addison-Wesley professional computing series. Addison-Wesley, Reading, Mass. [u.a.], 3rd print. edition, 1999.