

Operating Systems

13. Condition Variables and Semaphores

Prof. Dr. Frank Bellosa | WT 2020/2021

KARLSRUHE INSTITUTE OF TECHNOLOGY (KIT) – ITEC – OPERATING SYSTEMS



Concurrency Objectives


- **Mutal Exclusion** (e.g., thread A and B don't run at the same time)
 - solved with **locks (mutex)**
- **Ordering** (e.g., thread B runs after thread A did something)
 - solved with **condition variables** and **semaphores**
- **Condition Variable**: queue of waiting threads
 - B **waits** for a **signal on CV** before running **wait (CV, ...)**
 - A sends **signal to CV** when time for B to run **signal (CV, ...)**

Condition Variable

- `wait(cond_t *cv, mutex_t *lock)`
 - assumes the lock is held when `wait()` is called
 - puts caller to sleep + releases the lock (atomically)
 - when awoken, reacquires lock before returning
- `signal(cond_t *cv)`
 - wake a single waiting thread (if ≥ 1 thread is waiting)
 - if there is no waiting thread, just return, doing nothing
- `broadcast(cond_t *cv)`
 - wake all waiting threads (if ≥ 1 thread is waiting)
 - if there are no waiting thread, just return, doing nothing
- **Keep state in addition to CV's**
 - CV's are used to signal threads when state changes
 - If state is already as needed, thread doesn't wait for a signal!


Producer-Consumer Problem

- Consider the **producer-consumer problem** (also known as **bounded-buffer problem**)
 - A buffer is shared between a producer and a consumer (here: LIFO)
 - An integer **count** keeps track of the number of currently available (previously produced) items
 - Every time, the **producer** produces an item, it places it in the buffer and increments **count**
 - When the buffer is full, the producer needs to sleep until the consumer consumed an item
 - When the **consumer** consumes an item, it removes the item from the buffer and decrements **count**
 - When the buffer is empty, the consumer needs to sleep until the producer produces an item




Producer-Consumer Problem

- Consider the **producer-consumer problem** (also known as **bounded-buffer problem**)
 - A buffer is shared between a producer and a consumer (here: LIFO)
 - An integer **count** keeps track of the number of currently available (previously produced) items
 - Every time, the **producer** produces an item, it places it in the buffer and increments **count**
 - When the buffer is full, the producer needs to sleep until the consumer consumed an item
 - When the **consumer** consumes an item, it removes the item from the buffer and decrements **count**
 - When the buffer is empty, the consumer needs to sleep until the producer produces an item




Producer-Consumer Problem

- Consider the **producer-consumer problem** (also known as **bounded-buffer problem**)
 - A buffer is shared between a producer and a consumer (here: LIFO)
 - An integer **count** keeps track of the number of currently available (previously produced) items
 - Every time, the **producer** produces an item, it places it in the buffer and increments **count**
 - When the buffer is full, the producer needs to sleep until the consumer consumed an item
 - When the **consumer** consumes an item, it removes the item from the buffer and decrements **count**
 - When the buffer is empty, the consumer needs to sleep until the producer produces an item




Producer-Consumer Problem

- Consider the **producer-consumer problem** (also known as **bounded-buffer problem**)
 - A buffer is shared between a producer and a consumer (here: LIFO)
 - An integer **count** keeps track of the number of currently available (previously produced) items
 - Every time, the **producer** produces an item, it places it in the buffer and increments **count**
 - When the buffer is full, the producer needs to sleep until the consumer consumed an item
 - When the **consumer** consumes an item, it removes the item from the buffer and decrements **count**
 - When the buffer is empty, the consumer needs to sleep until the producer produces an item



Producer-Consumer Problem

- Consider the **producer-consumer problem** (also known as **bounded-buffer problem**)
 - A buffer is shared between a producer and a consumer (here: LIFO)
 - An integer **count** keeps track of the number of currently available (previously produced) items
 - Every time, the **producer** produces an item, it places it in the buffer and increments **count**
 - When the buffer is full, the producer needs to sleep until the consumer consumed an item
 - When the **consumer** consumes an item, it removes the item from the buffer and decrements **count**
 - When the buffer is empty, the consumer needs to sleep until the producer produces an item



Producer-Consumer Problem

```
void producer()
{
    Item newItem;

    for(;;) // ever
    {
        newItem = produce();

        if( count == MAX_ITEMS )
            sleep();

        insert( newItem );
        count++;

        if( count == 1 )
            wake_up( consumer );
    }
}
```

```
void consumer()
{
    Item item;

    for(;;) // ever
    {
        if( count == 0 )
            sleep();

        item = remove();
        count--;

        if( count == MAX_ITEMS - 1 )
            wake_up( producer );

        consume( item );
    }
}
```

■ Race condition on count

Non-Solution with mutex

```
void producer()
{
    Item newItem;

    for(;;) // ever
    {
        newItem = produce();

        if( count == MAX_ITEMS )
            sleep();
        mutex_lock( &lock );
        insert( newItem );
        count++;
        mutex_unlock( &lock );
        if( count == 1 )
            wake_up( consumer );
    }
}
```

```
void consumer()
{
    Item item;

    for(;;) // ever
    {
        if( count == 0 )
            sleep();
        mutex_lock( &lock );
        item = remove();
        count--;
        mutex_unlock( &lock );
        if( count == MAX_ITEMS - 1 )
            wake_up( producer );

        consume( item );
    }
}
```

■ if statements can still be racy

Another non-Solution with mutex

```
void producer()
{
    Item newItem;

    for(;;) // ever
    {
        newItem = produce();
        mutex_lock( &lock );
        if( count == MAX_ITEMS )
            sleep();

        insert( newItem );
        count++;
        if( count == 1 )
            wake_up( consumer );
        mutex_unlock( &lock );
    }
}
```

```
void consumer()
{
    Item item;

    for(;;) // ever
    {
        mutex_lock( &lock );
        if( count == 0 )
            sleep();

        item = remove();
        count--;
        if( count == MAX_ITEMS - 1 )
            wake_up( producer );
        mutex_unlock( &lock );
        consume( item );
    }
}
```

- One cannot work while the other sleeps with lock held (**deadlock**)

Final non-Solution with mutex

```
void producer()
{
    [...]
    for(;;) // ever
    {
        newItem = produce();
        mutex_lock( &lock );
        if( count == MAX_ITEMS ) {
            mutex_unlock( &lock );
            sleep();
            mutex_lock( &lock );
        }
        insert( newItem );
        count++;
        if( count == 1 )
            wake_up( consumer );
        mutex_unlock( &lock );
    }
}
```

```
void consumer()
{
    [...]
    for(;;) // ever
    {
        mutex_lock( &lock );
        if( count == 0 )
        {
            mutex_unlock( &lock );
            sleep();
            mutex_lock( &lock );
        }
        item = remove();
        count--;
        if( count == MAX_ITEMS - 1 )
            wake_up( producer );
        mutex_unlock( &lock );
        consume( item );
    }
}
```

■ Still racy and can cause **wakeup loss**

Solution with 2 Condition Variables

- Two condition variables: **empty** and **fill**

```
void producer()
{
    Item newItem;

    for(;;) // ever
    {
        newItem = produce();

        mutex_lock( &lock );
        while( count == MAX_ITEMS )
            cond_wait( &empty, &lock );

        insert( newItem );
        count++;
        cond_signal( &fill );
        mutex_unlock( &lock );
    }
}
```

```
void consumer()
{
    Item item;

    for(;;) // ever
    {
        mutex_lock( &lock );
        while( count == 0 )
            cond_wait( &fill, &lock );

        item = remove();
        count--;

        cond_signal( &empty );
        mutex_unlock( &lock );
        consume( item );
    }
}
```

Alloc-Free with cond_broadcast

```
void *allocate(int size) {  
    mutex_lock(&m);  
    while (bytesLeft < size)  
        cond_wait(&c, &m);  
    ...  
}
```

```
void free(void *ptr, int size) {  
    ...  
    cond_broadcast(&c)  
    ...  
}
```

Rules of Thumb for CV's

- **Keep state in addition to CV's**
- **Always do wait/signal with lock held**
- **Whenever thread wakes from waiting, recheck state**
 - Use “while” instead of “if”
 - Some implementations also have “spurious wakeups”
(may wake multiple waiting threads at signal or at any time)

Semaphore

- Introduce two syscalls that operate on data structure with integer element that we call **semaphore**
 - `sem_wait(&s)`: if $s > 0$: $s--$ and continue. Otherwise let caller sleep.
 - `sem_post(&s)`: if no thread is waiting: $s++$. Otherwise wake one up.
- Initialize s to the maximum number of threads that may enter the CS at any given time
 - `sem_init(&s, initval)` user cannot write value after initialization
 - `sem_wait` corresponds to `enter_critical_section()`
 - `sem_post` corresponds to `leave_critical_section()`
 - If you want to be specific about your semaphore allowing more than one thread in the CS, you can call it **counting semaphore**
- A semaphore that is initialized to 1 is called **binary semaphore**, **mutex semaphore** or just **mutex**
 - A mutex only admits one thread into the CS at a time

Condition Variables vs. Semaphores

- Condition variables have no state (other than waiting queue)
 - Programmer must track additional state
- Semaphors have state: track integer value
 - State cannot be directly accessed by user program, but state determines behavior of semaphore operations
 - Each semaphore is also associated with a wake-up queue
 - **Weak semaphores** Wake up a random waiting thread on **post**
 - **Strong semaphores** Wake up thread strictly in the order in which they started **waiting**

Equivalence Claim

- Locks can be built from semaphores
- Condition Variables can be built from semaphores
- Semaphores can be built from locks + condition variables

Build Lock from Semaphore

```
typedef struct __lock_t {
    sem_t sem;
} lock_t;

void lock_init(lock_t *lock) {
    sem_init(&lock->sem, 1); // 1 thread can grab lock
}

void lock_acquire(lock_t *lock) {
    sem_wait(&lock->sem);
}

void lock_release(lock_t *lock) {
    sem_post(&lock->sem);
}
```

Building CV's over Semaphores

- Possible, but really hard to do right
- Read about Microsoft Research's attempts:
<http://research.microsoft.com/pubs/64242/ImplementingCVs.pdf>

Build Semaphore from Lock and CV

```
typedef struct {
    int value;
    cond_t cond;
    lock_t lock;
} sem_t;

void sem_init(sem_t *s, int value) {
    s->value = value;
    cond_init(&s->cond);
    lock_init(&s->lock);
}
```

Build Semaphore from Lock and CV

```
sem_wait(sem_t *s) {  
    lock_acquire(&s->lock);  
  
    while (s->value <= 0)  
        cond_wait(&s->cond, &s->lock);  
  
    s->value--;  
  
    lock_release(&s->lock);  
}
```

```
sem_post(sem_t *s) {  
    lock_acquire(&s->lock);  
  
    s->value++;  
  
    cond_signal(&s->cond);  
  
    lock_release(&s->lock);  
}
```

Bounded Buffer with Semaphores

```
sem_t emptyBuffer, fullBuffer; // counting semaphores
sem_t mutex; // binary semaphore

init() {
    sem_init(&emptyBuffer, ITEMS); // all slots empty
    sem_init(&fullBuffer, 0);      // no filled slots
    sem_init(&mutex, 1);           // mutex
}
```

Bounded Buffer with Semaphores

```
void producer() {  
  
    for(;;) { // for ever  
        sem_wait(&emptyBuffer);  
        // wait for free slot  
        sem_wait(&mutex);  
        slot = findempty(&buffer);  
        sem_signal(&mutex);  
  
        fill(&buffer[slot]);  
        sem_signal(&fullBuffer);  
        // signal filled slot  
    }  
}
```

```
void consumer() {  
  
    for(;;) { // for ever  
        sem_wait(&fullBuffer);  
        // wait for filled slot  
        sem_wait(&mutex);  
        slot = findfull(&buffer);  
        sem_signal(&mutex);  
  
        use(&buffer[slot]);  
        sem_signal(&emptyBuffer);  
        // signal use of slot  
    }  
}
```


Readers-Writers Problem

- Problem: Model access to shared data structures
 - Many threads compete to read or write the same data
 - **Readers** only read the data set; they do not perform any updates
 - **Writers** can both read and write
- Using a single mutex for read and write operations is not a good solution, as it unnecessarily blocks out multiple readers while no writer is present
- Idea: Locking should reflect different semantics for reading data and for writing data
 - If no thread writes, multiple readers may be present
 - If a thread writes, no other readers and writers are allowed

1st Readers-Writers Problem: Readers Preference

- No reader should have to wait if other readers are already present

```
typedef struct _rwlock_t{
    int reader;
    sem_t lock;
    sem_t writelock;
} rwlock_t;

void rwlock_init(rwlock_t *rw) {
    rw->readers = 0;
    sem_init(&rw->lock, 1);
    sem_init(&rw->writelock, 1);
}
```

1st Readers-Writers Problem: Readers Preference

- No reader should have to wait if other readers are already present

```
void writer() {
    for(;;) { // ever
        // generate data to write
        sem_wait(&rw->writelock);
        // write data
        sem_post(&rw->writelock);
    }
}
```

- Writers cannot acquire **writelock** until the last reader leaves the critical section

```
void reader() {
    for(;;) { // ever
        sem_wait(&rw->lock);
        rw->readers++;
        if (rw->readers == 1)
            sem_wait(&rw->writelock);
        sem_post(&rw->lock);
        // read data
        sem_wait(&rw->lock);
        rw->readers--;
        if (rw->readers == 0)
            sem_post(&rw->writelock);
        sem_post(&rw->lock);
    }
}
```

2nd Readers-Writers Problem: Writers Preference

- No writer shall be kept waiting longer than absolutely necessary
- Code is analogous to 1st readers-writers problem but with separate readers- and writers-counts
- Read “Concurrent Control with Readers and Writers” [CHP71]
- 1st and 2nd readers-writers problem have the same issue:
 - Readers preference → writers can starve
 - Writers preference → readers can starve

3rd Readers-Writers Problem: Bounded Waiting

- No thread shall starve
- POSIX threads contains [readers-writers locks](#) to address this issue
- Multiple readers but only a single writer are let into the CS
- If readers are present while a writer tries to enter the CS then
 - don't let further readers in
 - block until readers finish
 - let writer in

References I

- [CHP71] P. J. Courtois, F. Heymans, and D. L. Parnas. Concurrent control with “readers” and “writers”. *Commun. ACM*, 14(10):667–668, October 1971.